

Colocation Explained

Heartbeat 2.1.2-4 Onwards

abeekhof@suse.de

Terminology



- ✦ `Collocate(B, A)`
- ✦ `<rsccollocation from=B to=A/>`
- ✦ Decide where to put A, then put B there too
- ✦ Include B's preferences when deciding where to put A
- ✦ If A cannot run anywhere, B can't run either
- ✦ If B cannot run anywhere, A will be unaffected

Adding Scores

- ✧ $\text{number} > \text{INFINITY} = \text{INFINITY}$
- ✧ $\text{number} < -\text{INFINITY} = -\text{INFINITY}$
- ✧ $\text{number} + \text{INFINITY} = \text{INFINITY}$
- ✧ $\text{number} - \text{INFINITY} = -\text{INFINITY}$
- ✧ $\text{INFINITY} - \text{INFINITY} = -\text{INFINITY}$

Simple Example

Setup

- ✦ resource(A, priority=5)
- ✦ resource(B, priority=50)
- ✦ location(A, node1, 100)
- ✦ location(A, node2, 10)
- ✦ location(B, node2, 1000)
- ✦ colocate(B, A)

Simple Example

What Happens

- ✦ Start at highest priority resource (**B**)
- ✦ Defer and process **A** instead (collocation rule)
- ✦ Incorporate **B**'s preferences
 - ✦ $A.\text{node1}.\text{score} += B.\text{node1}.\text{score}$ (100)
 - ✦ $A.\text{node2}.\text{score} += B.\text{node2}.\text{score}$ (1010)
- ✦ Choose a node (**node2**)

Simple Example

Actually I Lied

- ✧ Incorporate **B**'s preferences
 - ✧ $A.\text{node}[x].\text{score} += \textbf{factor} * B.\text{node}[x].\text{score}$
- ✧ What is ***factor***?
 - ✧ $\text{factor} ::= \text{constraint.score} / \text{INFINITY}$
- ✧ *For most people it will be 1 or -1*
- ✧ *So really its: $\text{colocate}(B, A, \textbf{score})$*

Choosing a Node for **B**

Simple Example

- ✧ Process collocation constraint
 - ✧ Matching node: $\text{node.score} = \text{INFINITY}$
 - ✧ Everything else: $\text{node.score} = -\text{INFINITY}$
- ✧ Scores do **not** include **A**'s preferences
- ✧ Final scores for **B**
 - ✧ $\text{node1} = -\text{INFINITY}$
 - ✧ $\text{node2} = \text{INFINITY}$

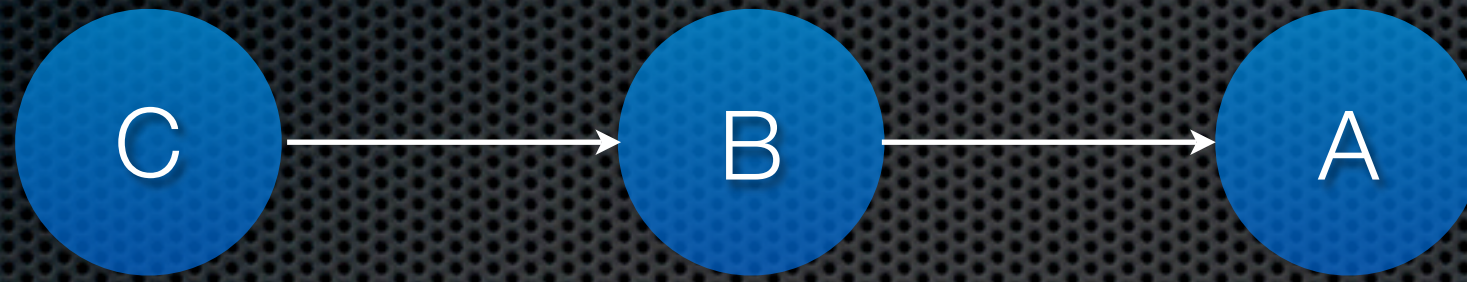
Choosing a Node for **B**

Suggested Collocation

- ✧ When the collocation score \neq INFINITY
 - ✧ Matching node: $\text{node.score} += \text{collocation.score}$
 - ✧ Everything else: unchanged
- ✧ Scores do **not** include **A**'s preferences
- ✧ Final scores for **B** (collocation.score = 500)
 - ✧ $\text{node1} = 0$
 - ✧ $\text{node2} = 1500$

Chained Example

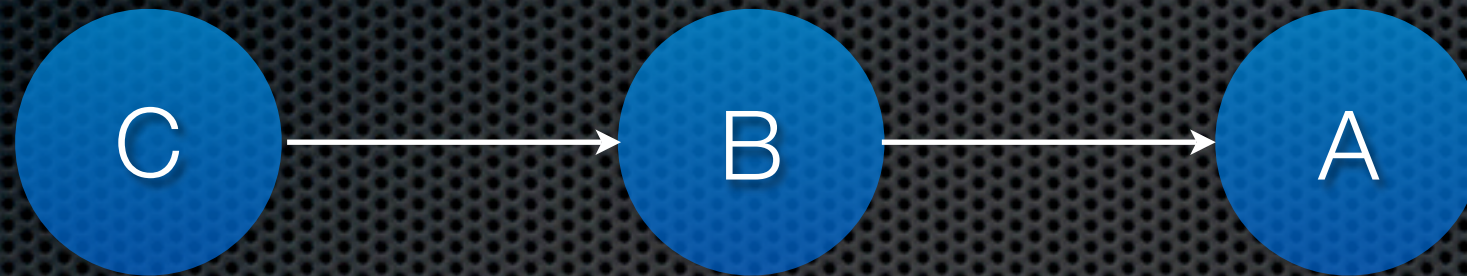
Setup



- ✦ resource(A, p=5)
- ✦ resource(B, p=500)
- ✦ resource(C, p=50)
- ✦ location(A, node1, 100)
- ✦ location(A, node2, 10)
- ✦ location(B, node2, 1000)
- ✦ location(C, node1, 10000)
- ✦ colocate(B, A)
- ✦ colocate(C, B)

Chained Example

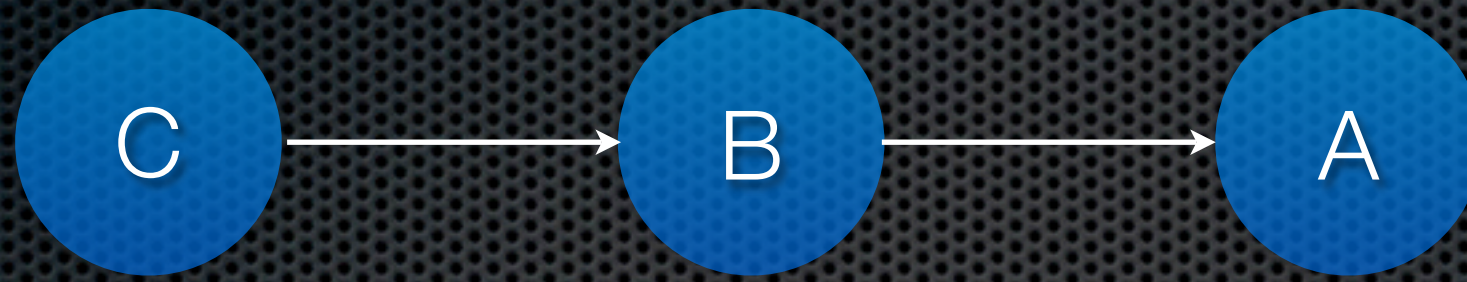
What Happens



- ✦ Start at highest priority resource (**B**)
- ✦ Defer and process **A** instead (collocation rule)
- ✦ Incorporate **B**'s preferences
 - ✦ $A.\text{node}[x].\text{score} += B.\text{node}[x].\text{score}$
- ✦ So far nothing is different

Chained Example

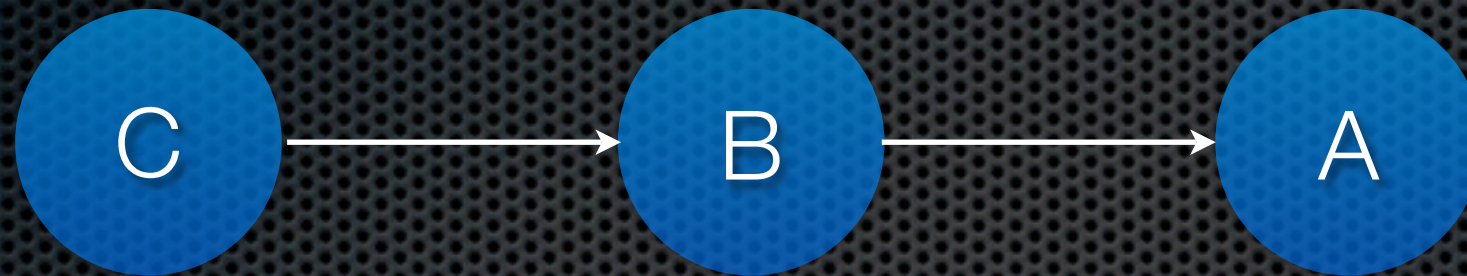
What Happens (Continued)



- ✦ Incorporate **C**'s preferences too!
 - ✦ $A.\text{node}[x].\text{score} += C.\text{node}[x].\text{score}$
- ✦ Final scores (when choosing a node for **A**)
 - ✦ $\text{node1} = 10100$
 - ✦ $\text{node2} = 1010$

Chained Example

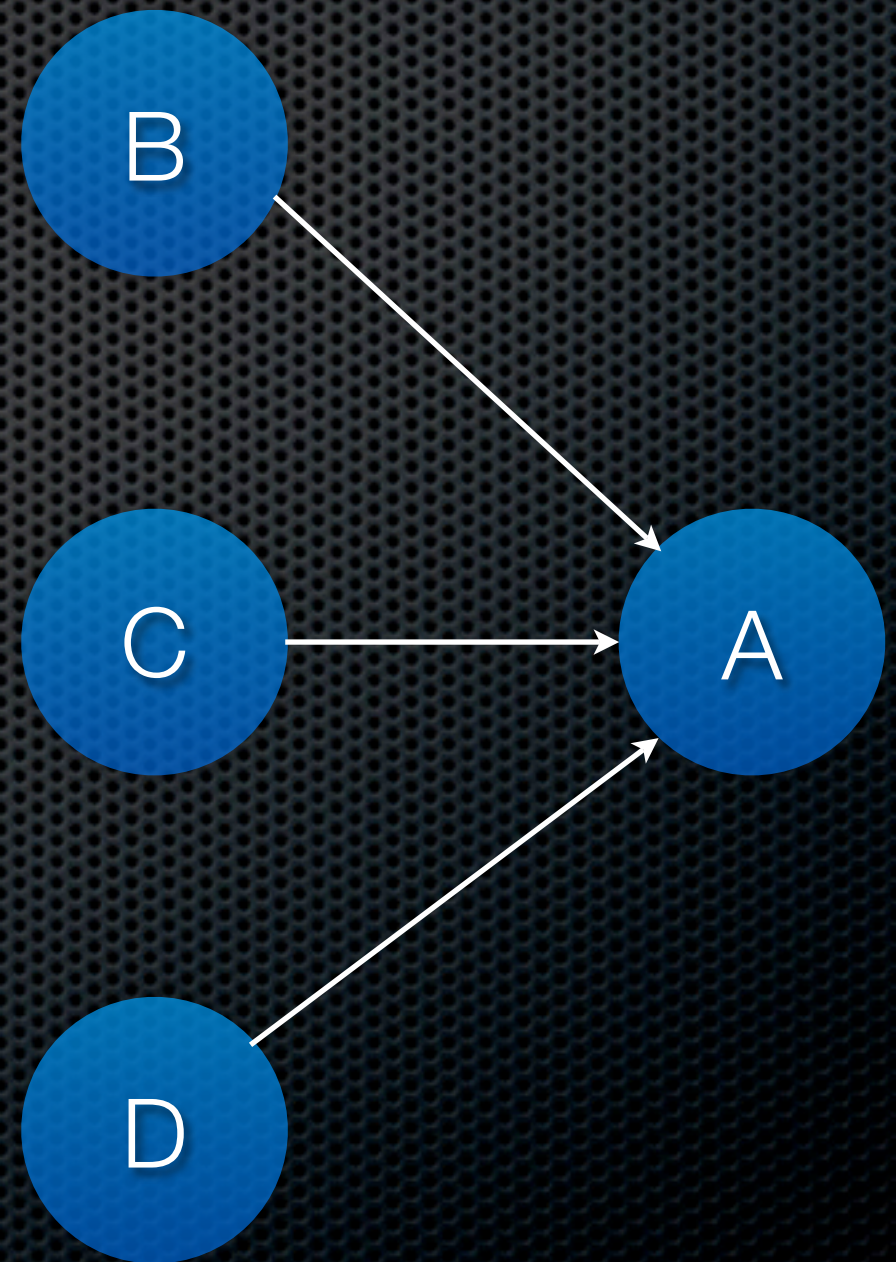
Final Scores: **B** and **C**



- ✦ Resource **B**
 - ✦ node1 = INFINITY
 - ✦ node2 = -INFINITY
- ✦ Resource **C**
 - ✦ node1 = INFINITY
 - ✦ node2 = -INFINITY

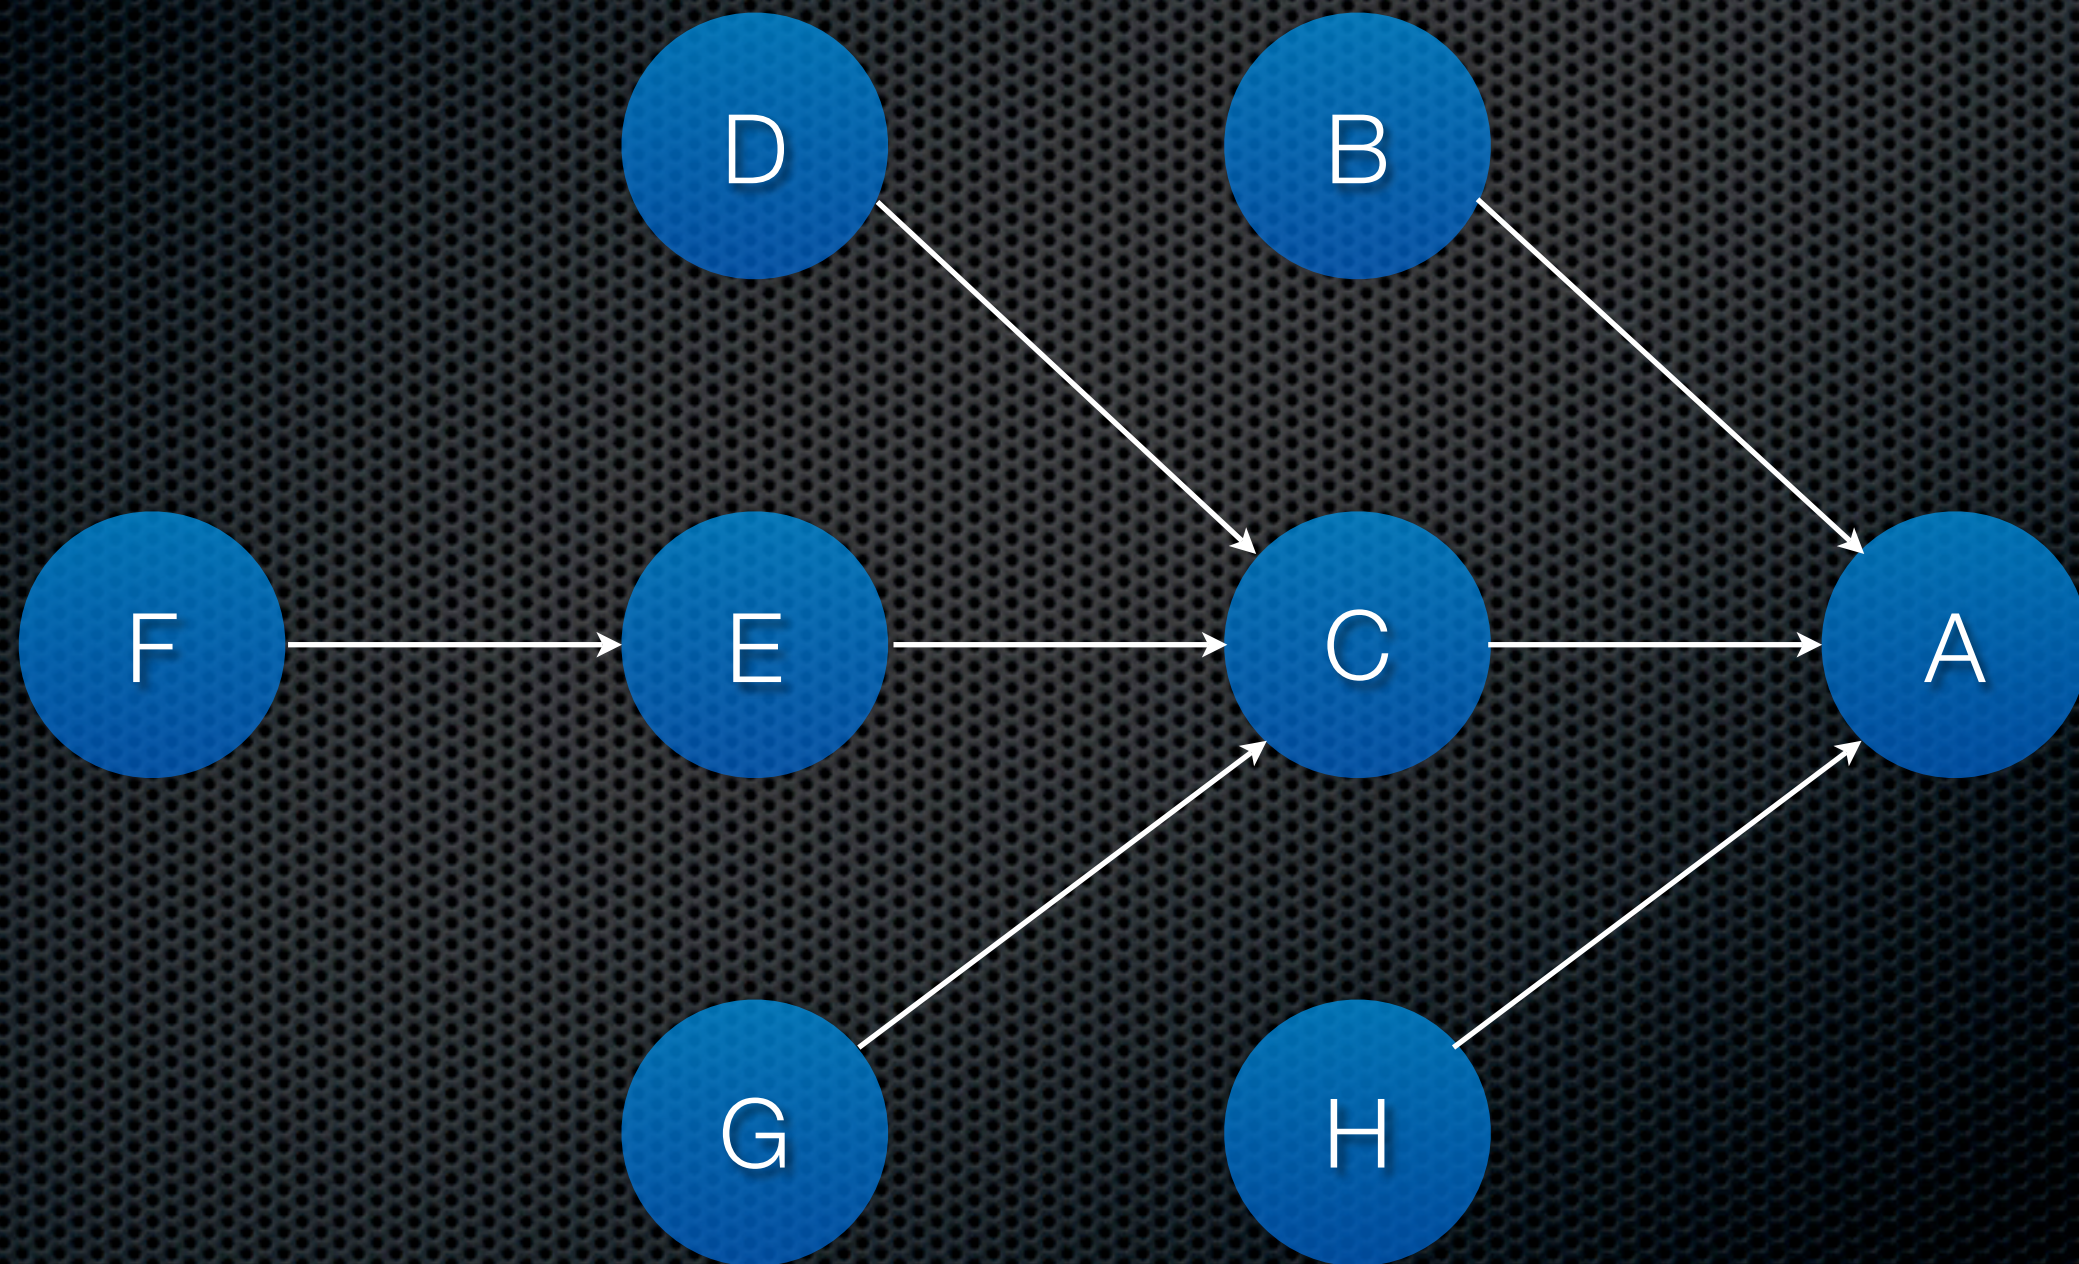
Multiple Dependencies

- ✦ Include scores from B, C and D when choosing a node for A
- ✦ Order is defined by priority of dependent resources (or name if priority is equal)
- ✦ In this example:
 - ✦ $B.\text{priority} > C.\text{priority}$
 - ✦ $C.\text{priority} > D.\text{priority}$



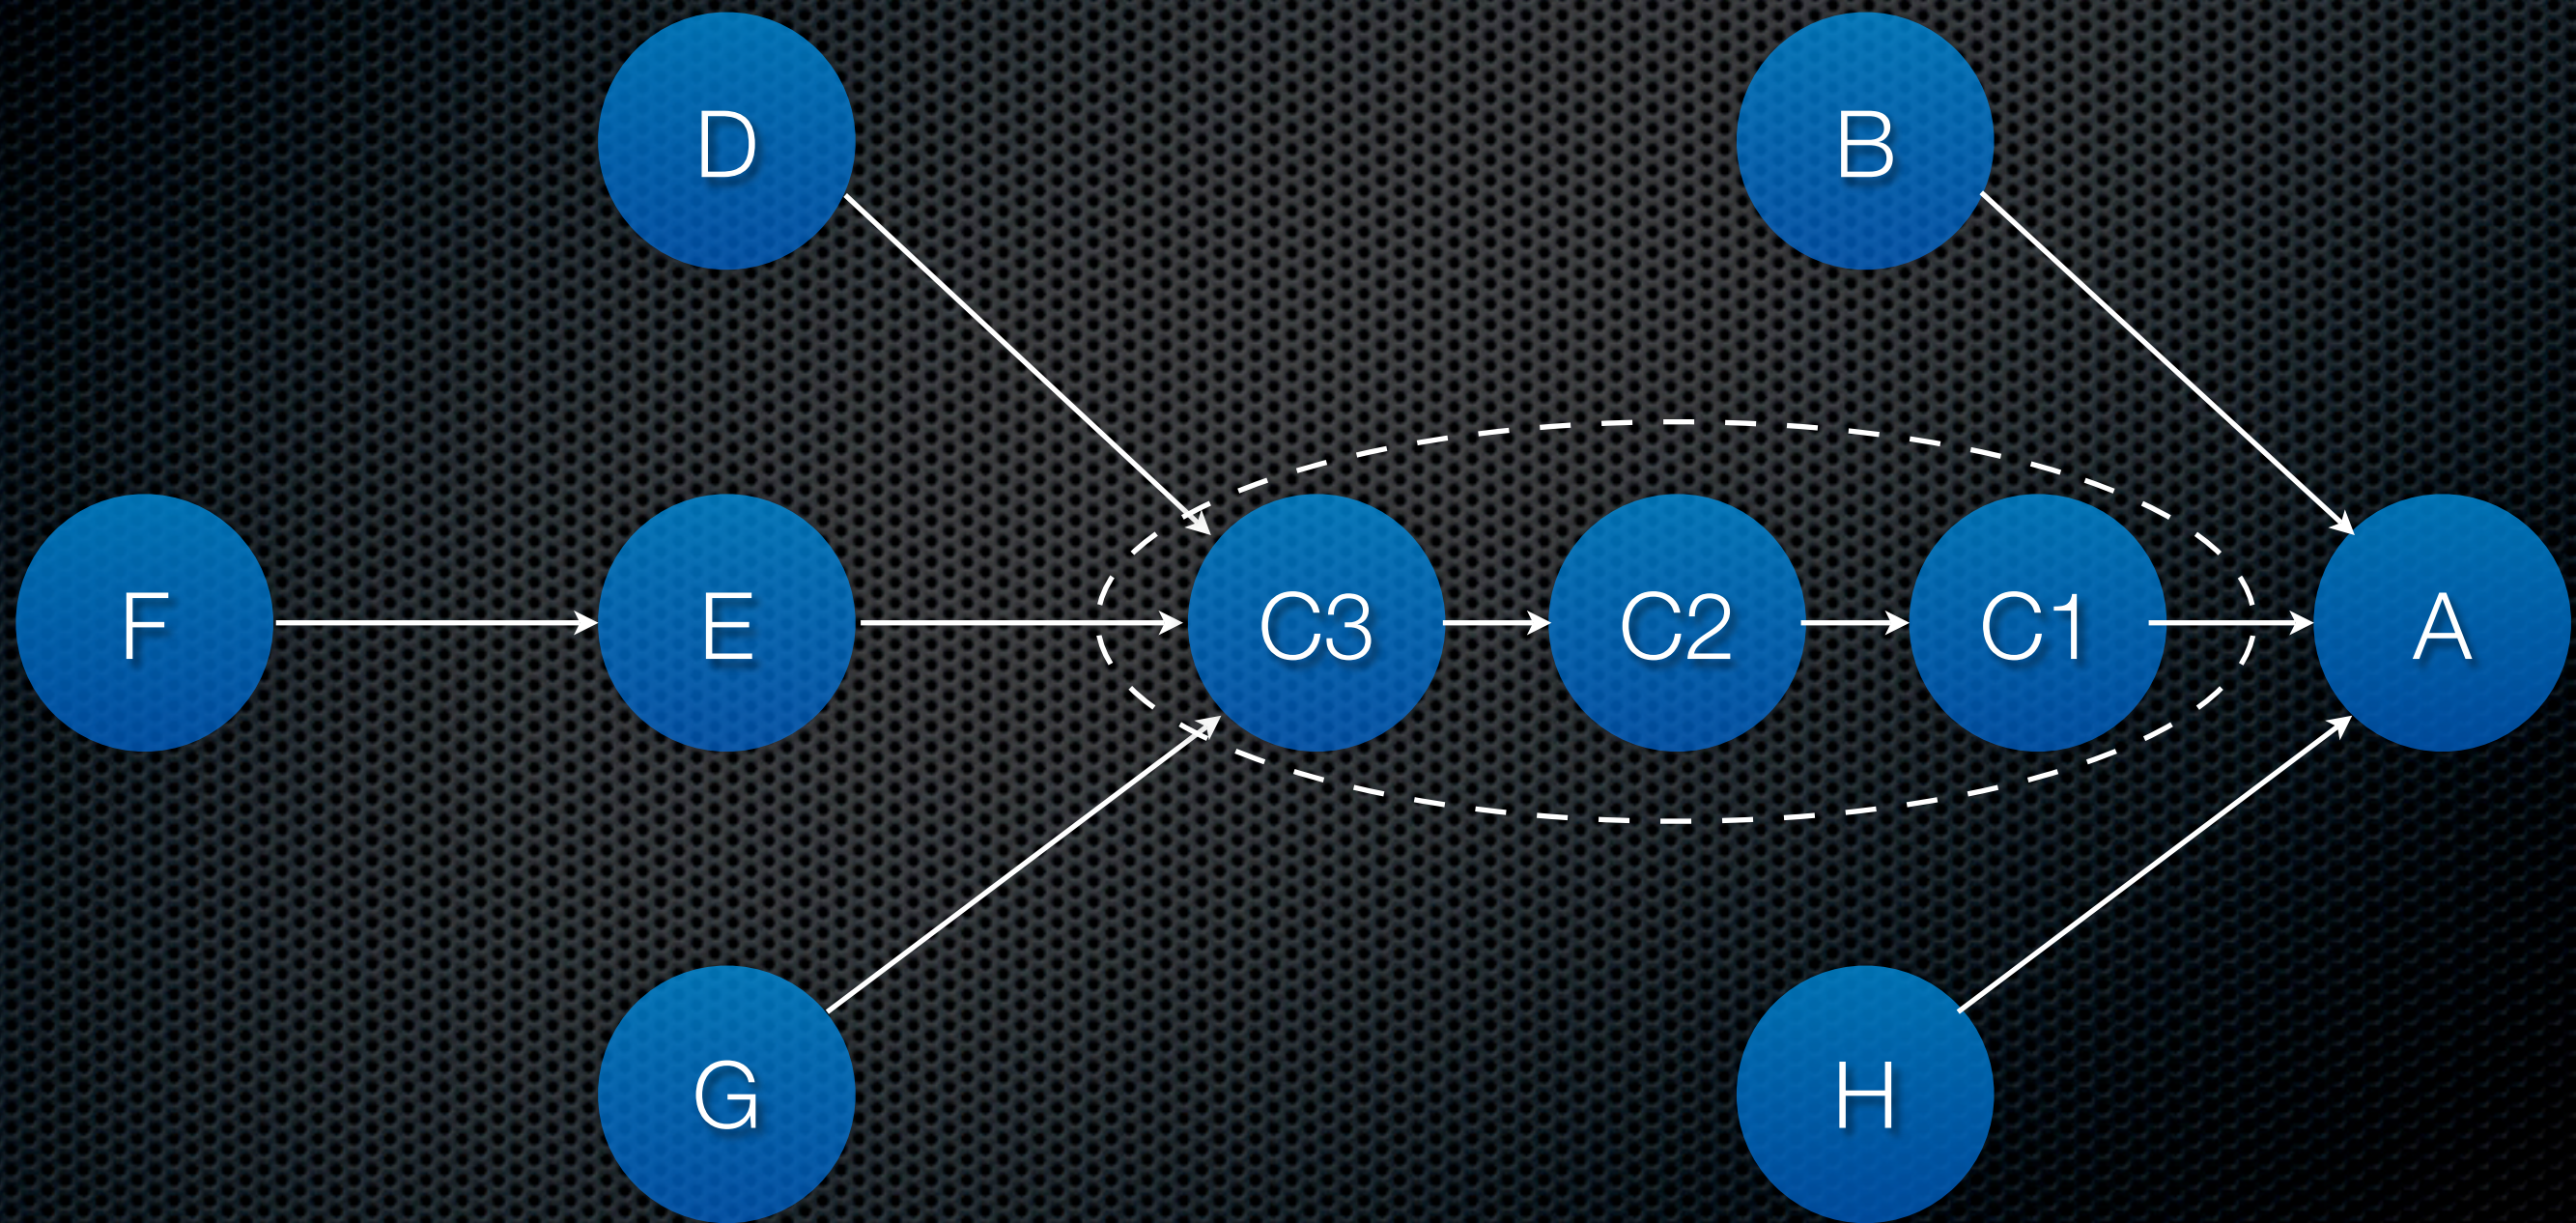
Dependancy Tree

Order in Which Preferences are Applied (A-H)



More Complex

C is a Group

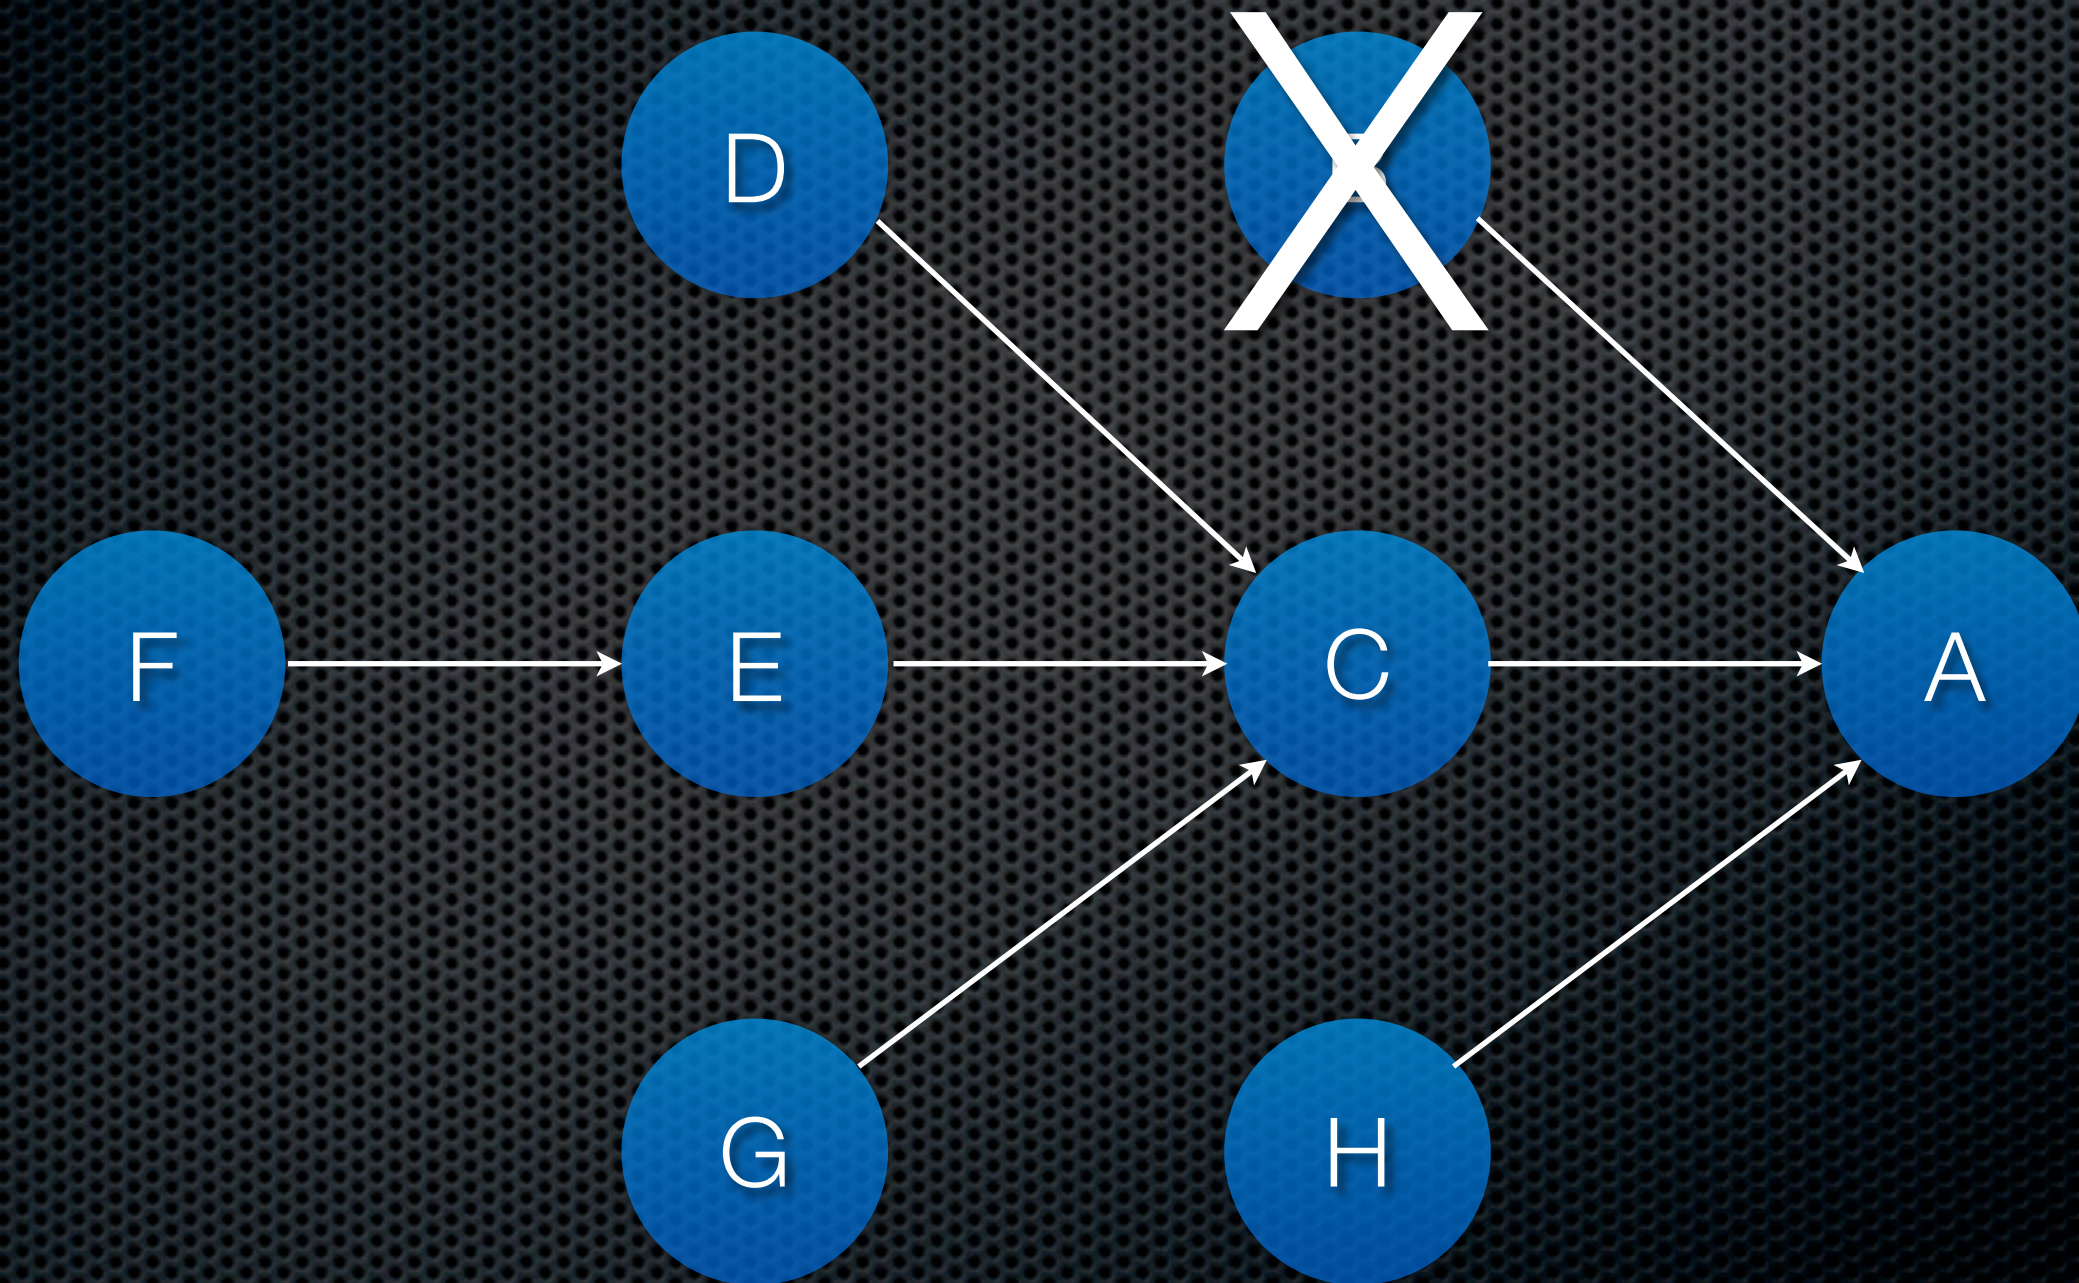


Getting Smart

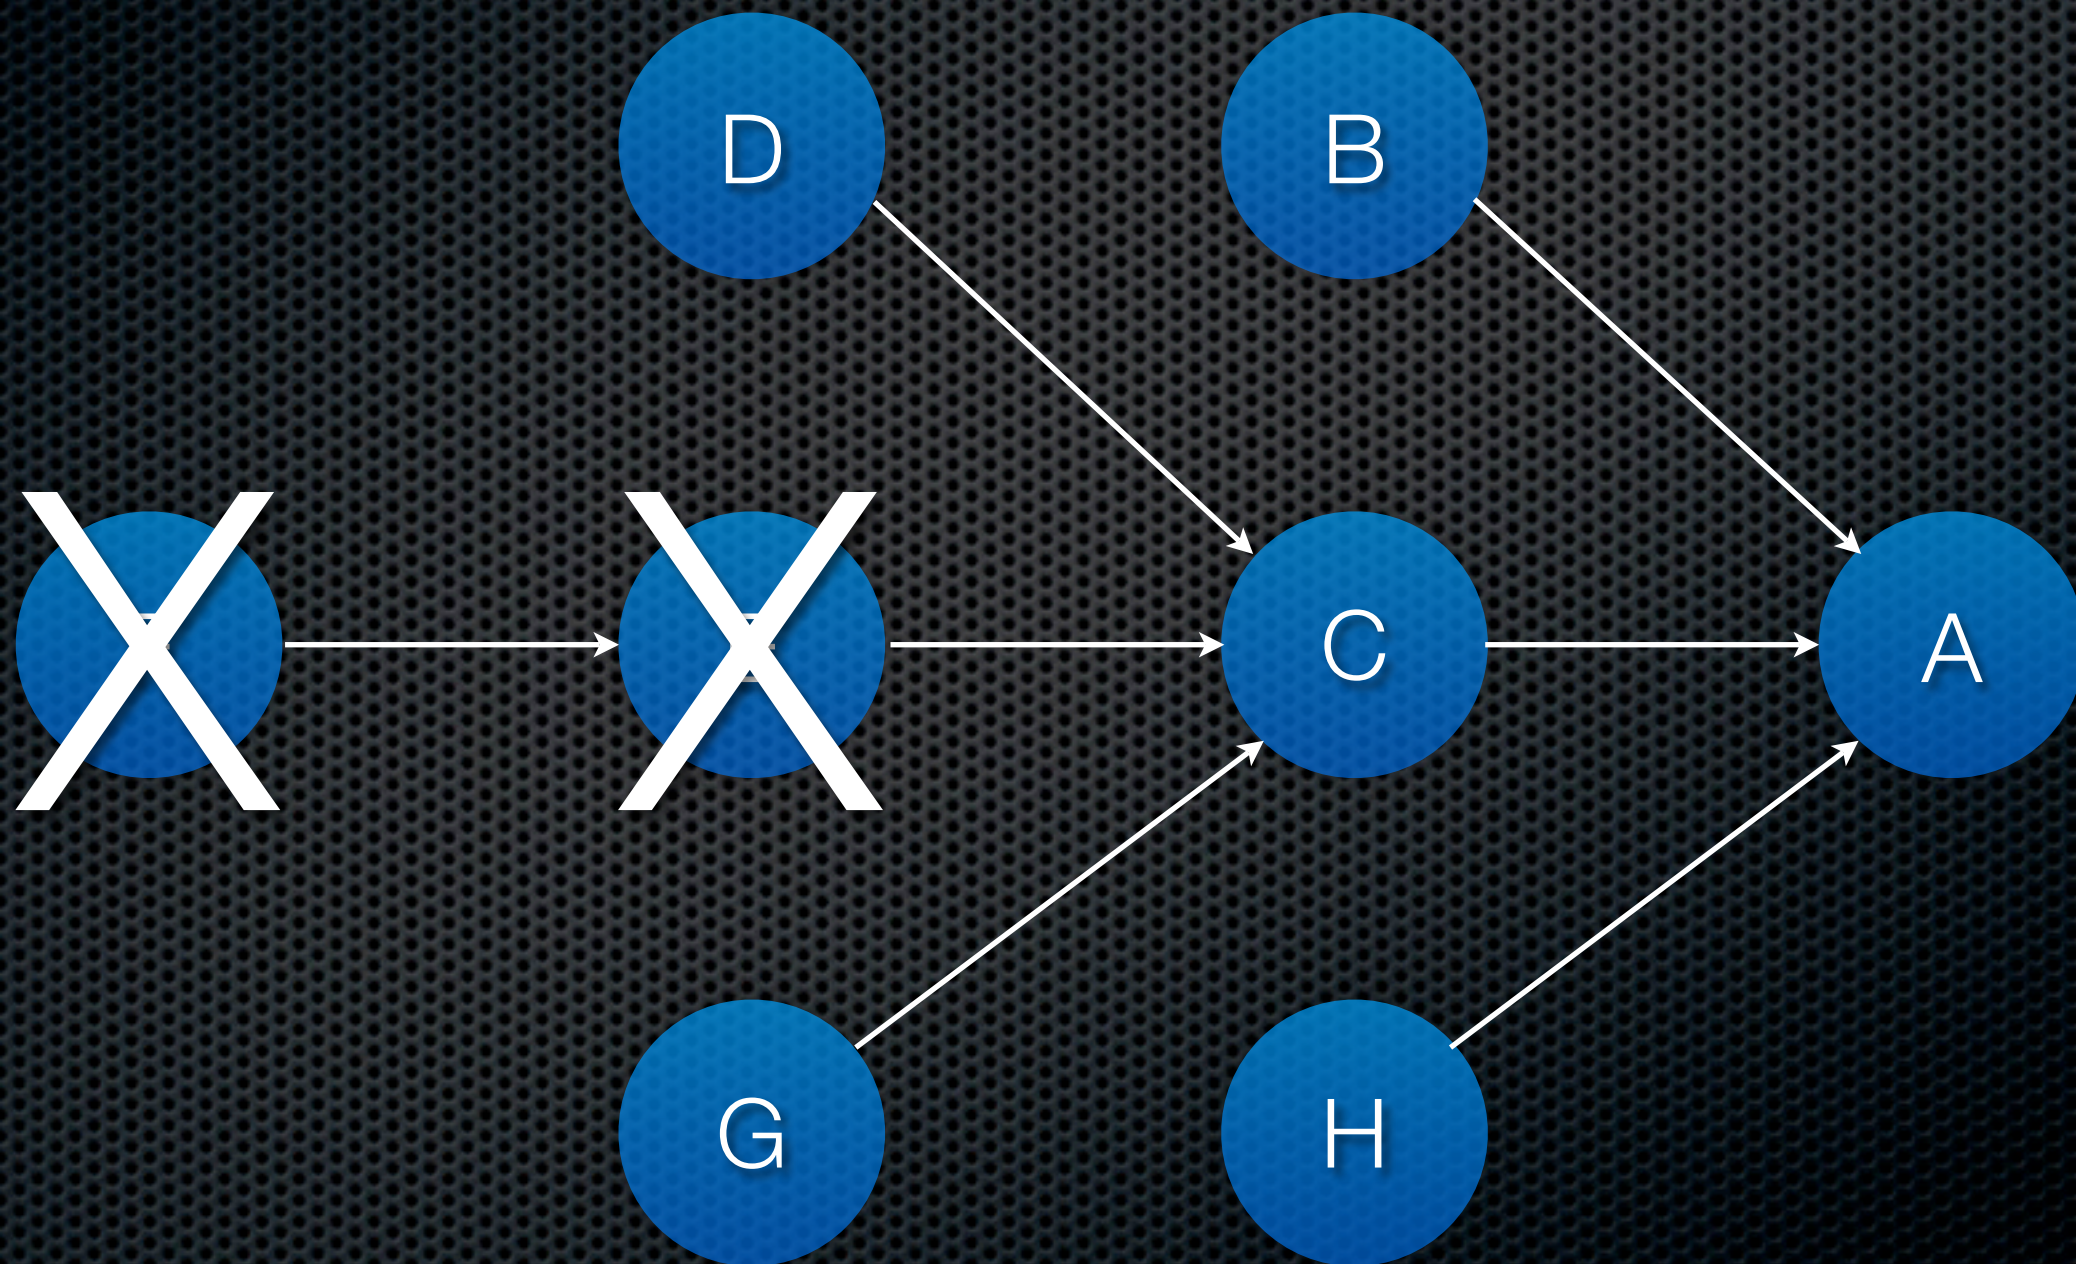
When not Everything can Run

- ✦ If applying a resource's preference, means that all nodes would be unavailable...
 - ✦ Undo the current resource's preference
 - ✦ Skip any resources that need to be colocated with the current resource
 - ✦ Process the next peer

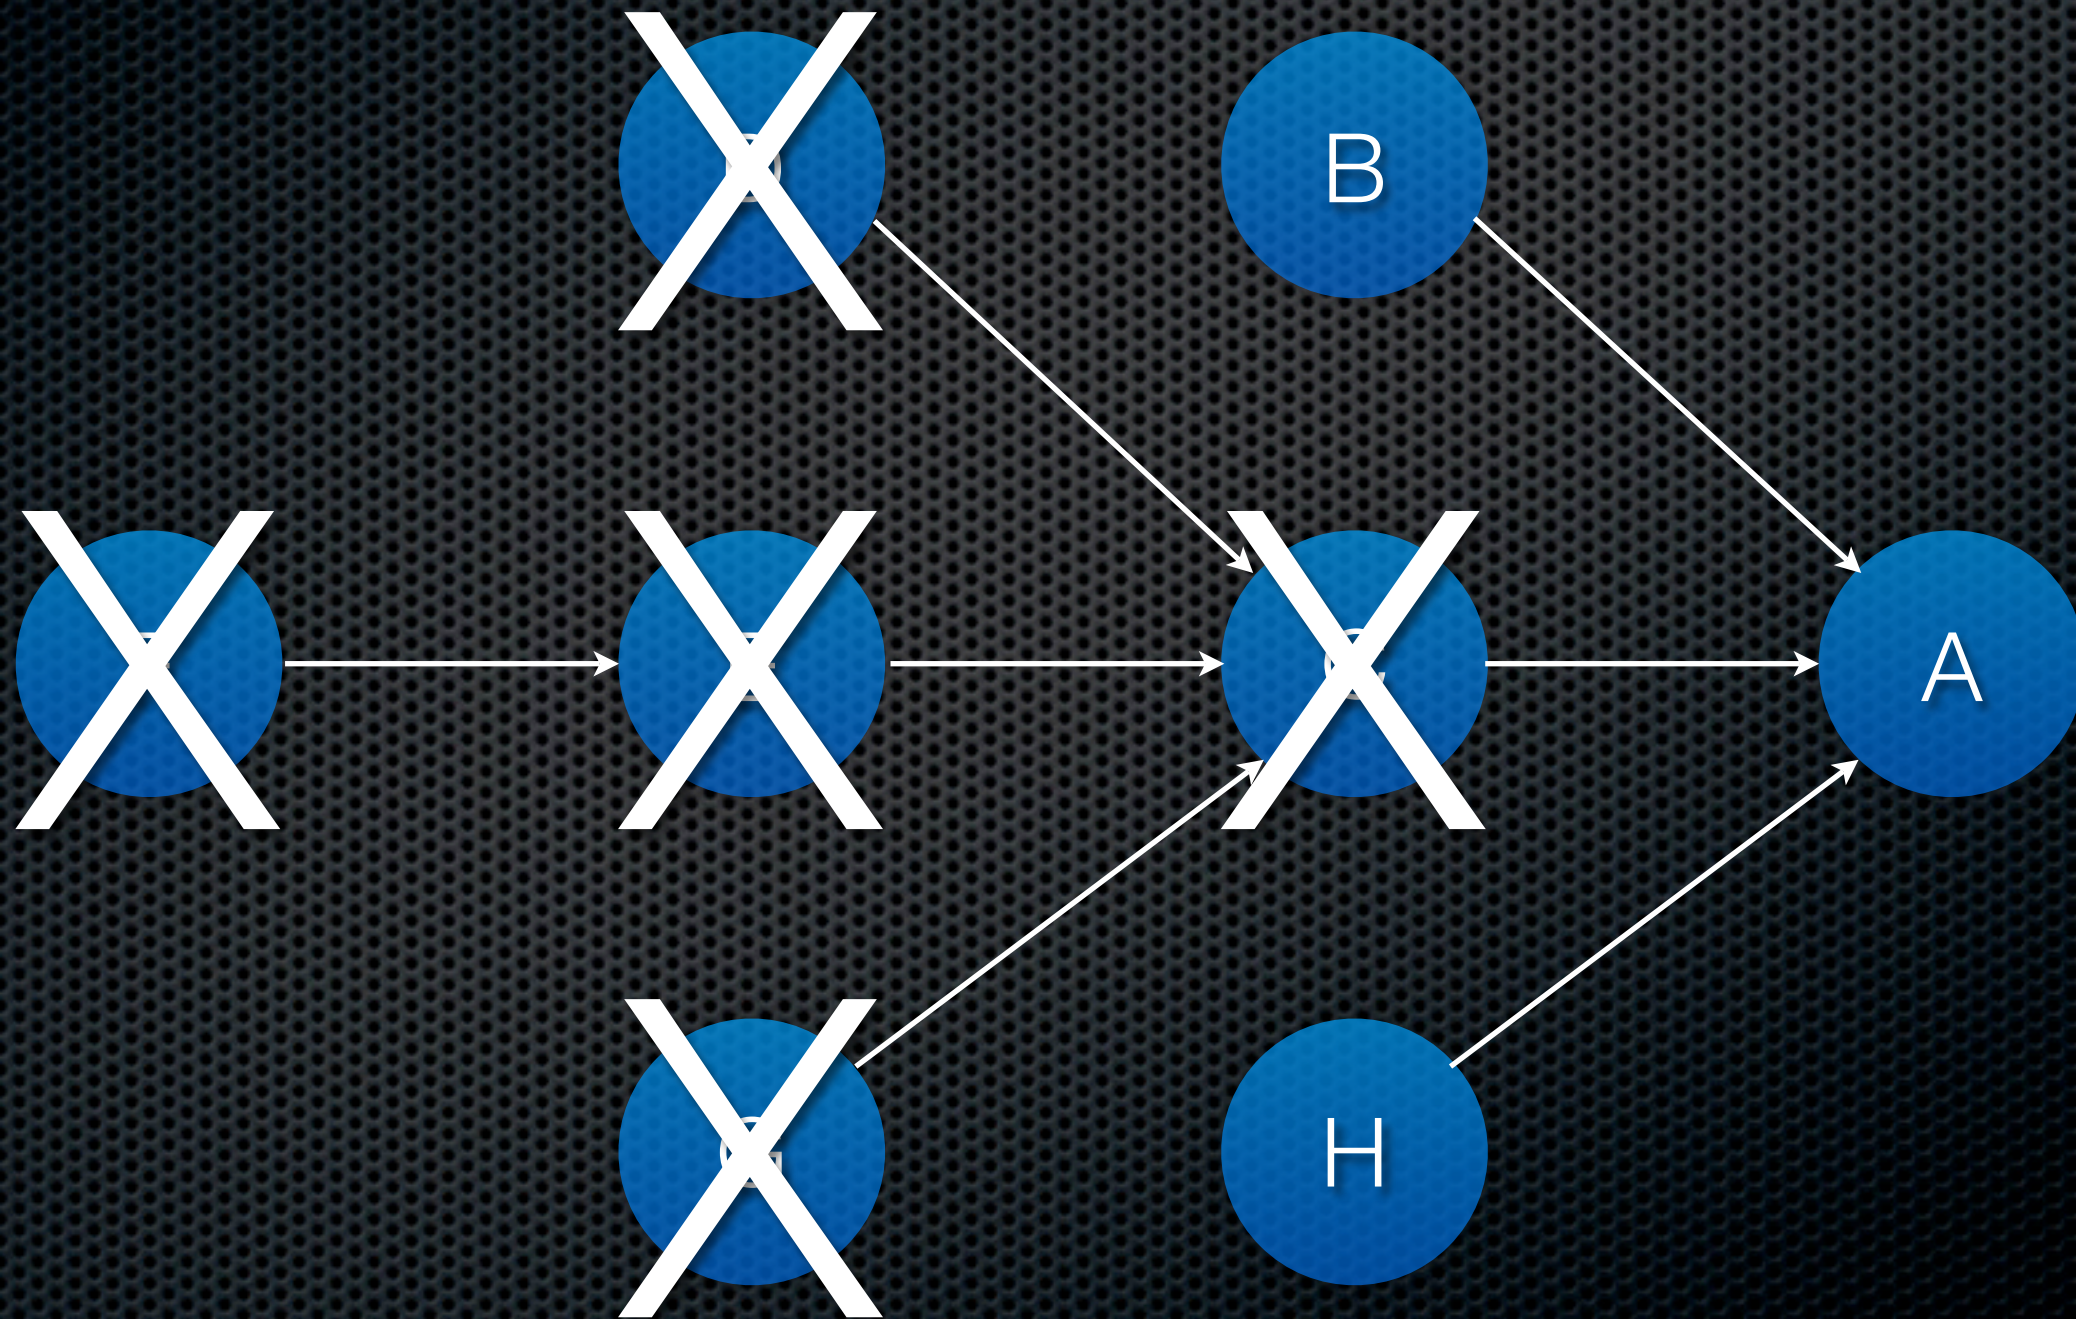
Un-runnable: B



Un-runnable: E



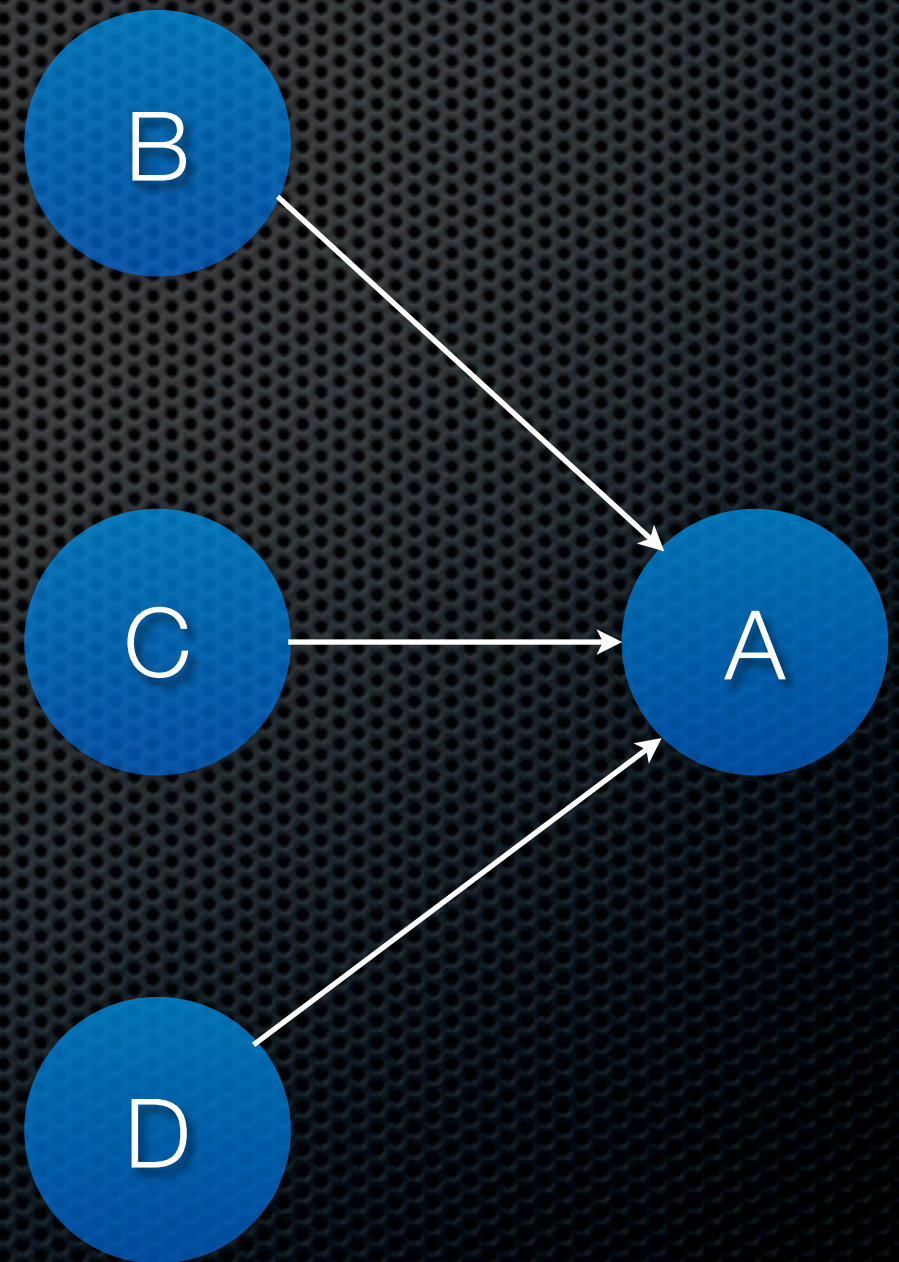
Un-runnable: C



Un-runnable

Worked Example

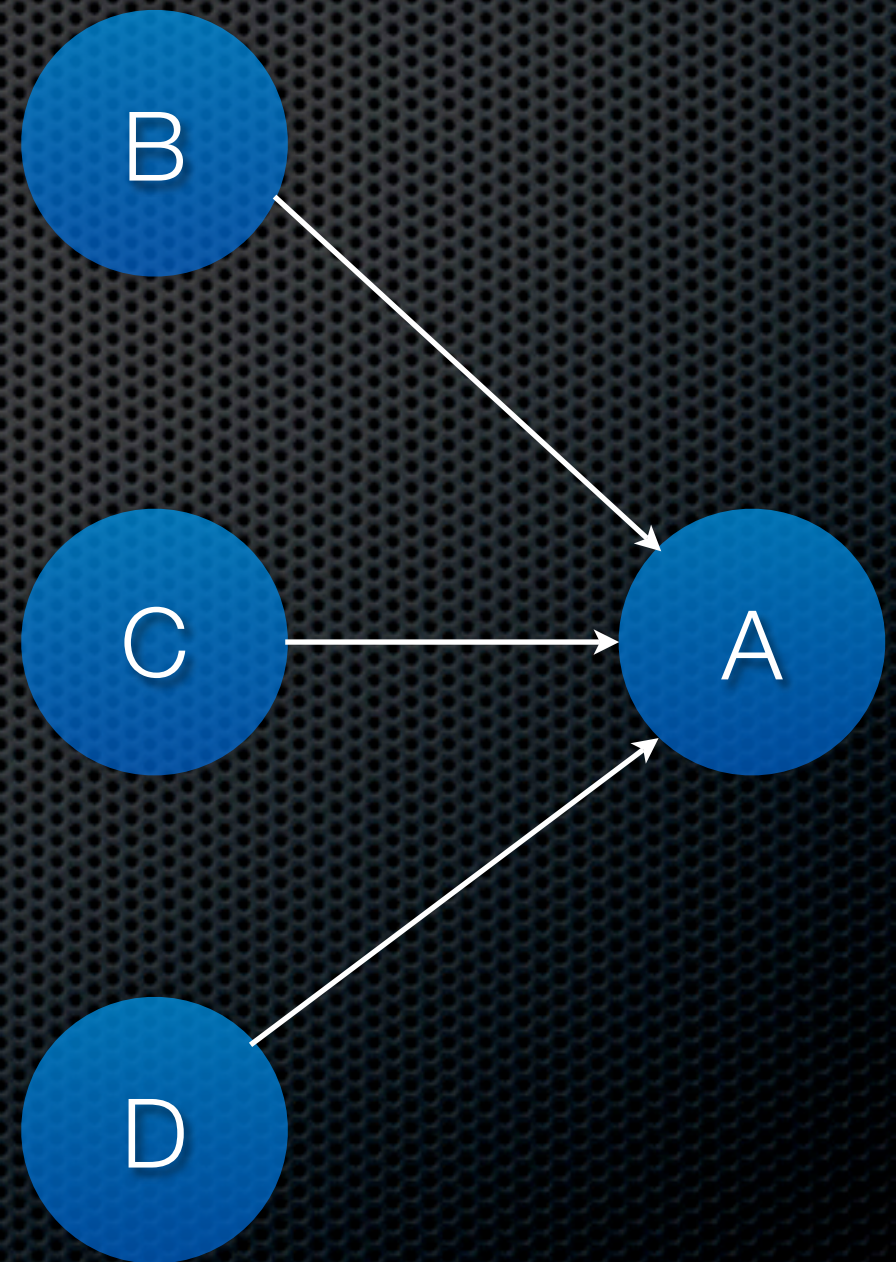
Rsc	Node	Score
A	node1	50
A	node2	5
B	node1	1
B	node2	10
C	node1	-INFINITY
C	node2	-INFINITY
D	node1	100



Un-runnable

Worked Example (continued)

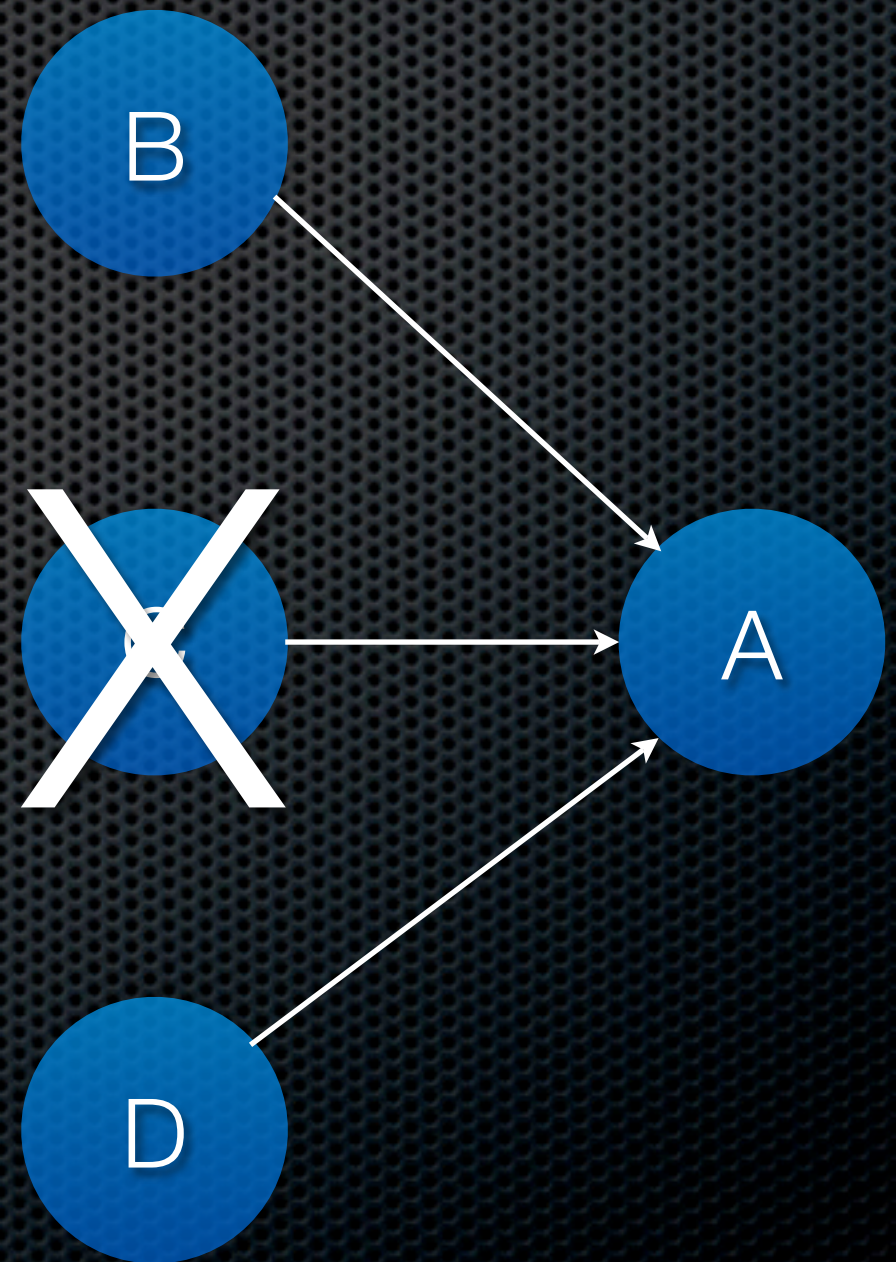
- ✧ Consider **B**
 - ✧ $A.\text{node1}.\text{score} = 50 + 1$
 - ✧ $A.\text{node2}.\text{score} = 5 + 10$



Un-runnable

Worked Example (continued)

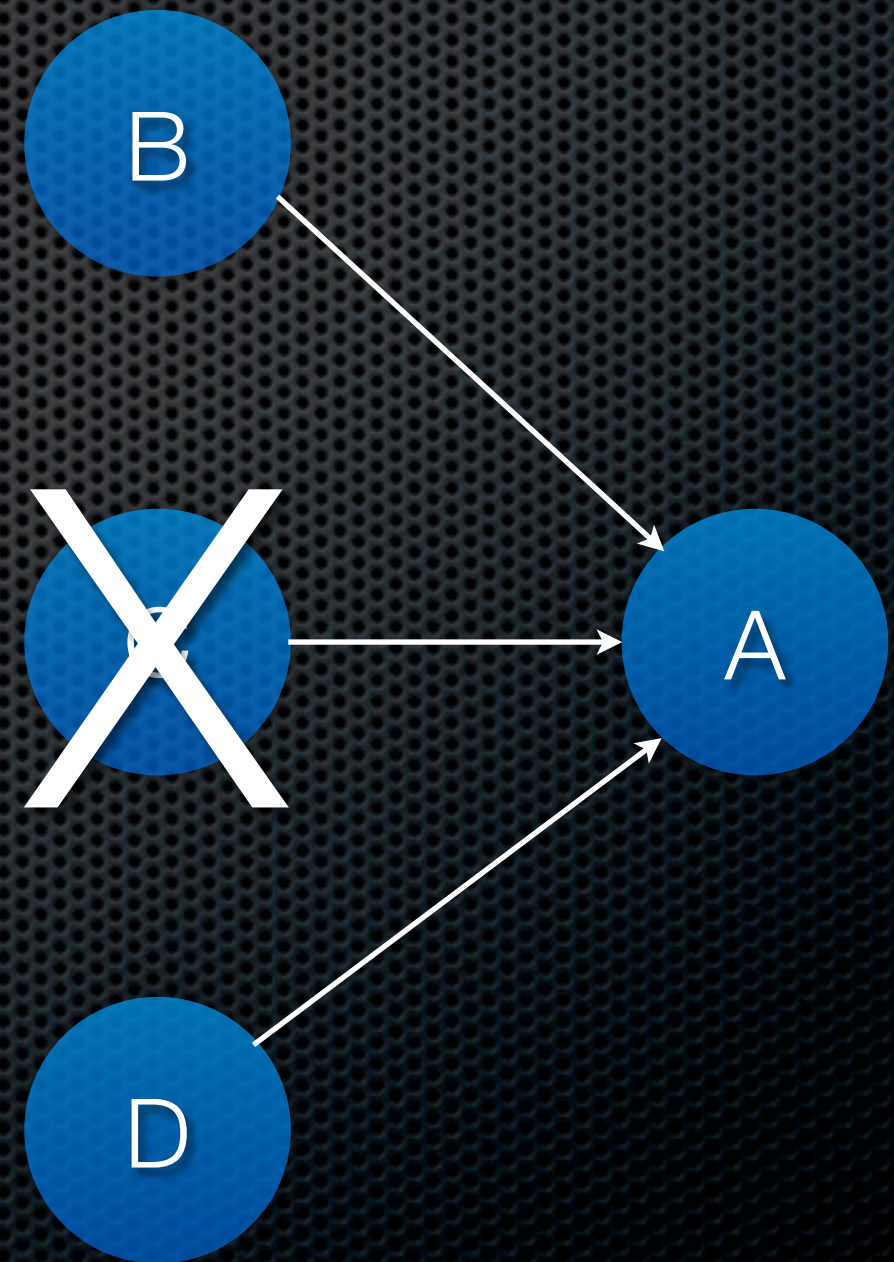
- ✦ Consider **C**
 - ✦ $A.\text{node1}.\text{score} = 51 - \text{INFINITY}$
 - ✦ $A.\text{node2}.\text{score} = 15 - \text{INFINITY}$
- ✦ Rollback Scores
 - ✦ $A.\text{node1}.\text{score} = 51$
 - ✦ $A.\text{node2}.\text{score} = 15$



Un-runnable

Worked Example (continued)

- ✧ Consider **D**
 - ✧ $A.\text{node1}.\text{score} = 51 + 100$
 - ✧ $A.\text{node2}.\text{score} = 15 + 1000$
- ✧ Final Scores
 - ✧ $A.\text{node1}.\text{score} = 151$
 - ✧ $A.\text{node2}.\text{score} = 1015$
- ✧ Choose **node2**



Colocation by Role

Master/Slave - Summary

- A resource that needs to run on the master can force the master to move (rather than not be allowed to run anywhere)
- A resource that can't run anywhere and must run with the master does not prevent the promotion of a master

Colocation by Role

Who Gets Promoted

- ✦ Allocation occurs as-per previous slides
- ✦ Decision of which instances to promote is based on
 - ✦ Preference as set by RA with `crm_master`
 - ✦ **Location preferences of resources that wish to be colocated with the master instance(s)**

Colocation by Role

Master/Slave Example

Child	Location	M/S Score
ms:0	node1	1000
ms:1	node2	100
ms:2	node3	10
ms:3	node4	-INFINITY

Colocation by Role

Changes

- ✦ Under the old system, we would
 - ✦ sort the children by their m/s score
 - ✦ allocate masters in that order (ms:0, ms:1, ms:2)
- ✦ Now we include the colocation scores too

Colocation by Role

Master/Slave Example (continued)

Dependent	Location	Score
rsc1	node1	20
rsc2	node2	200
rsc3	node2	-INFINITY
rsc3	node3	2000
rsc4	[everywhere]	-INFINITY

Colocation by Role

Master/Slave Example (continued)

Child	Location	M/S Score	Final Score
ms:0	node1	1000	1020
ms:1	node2	100	-INFINITY
ms:2	node3	10	2010
ms:3	node4	-INFINITY	-INFINITY

Colocation by Role

Master/Slave Example (continued)

- ✦ “Final” weight affects sorting order only
 - ✦ Negative final score does not prevent the instance from being promoted
- ✦ Sort and allocate Masters in order (depending on the number of masters required):
 - ✦ ms:2 , ms:0, ms:1
- ✦ ms:3 can't be promoted as it's m/s score is less than zero



abeekhof@suse.de or linux-ha@lists.linux-ha.org